

# Abstract

In this tutorial, I analyze the critical relationship between a Multilayer Perceptron's (MLP) **Depth** (number of layers) and **Width** (number of neurons per layer) and its classification performance. I demonstrate how these architectural choices determine a network's **capacity** and directly influence the essential **Bias–Variance Trade-off**. Through exhaustive code examples, interactive tools, and visual analysis, I provide a practical, evidence-based guide for selecting the optimal MLP architecture—the "Goldilocks Zone"—to achieve maximum generalization.

## 1. Introduction: The Architecture of Learning

Neural Networks are powerful function approximators whose performance is highly dependent on their internal structure. For the Multilayer Perceptron (MLP), this structure is defined by its **architecture**:

- **Depth:** The number of hidden layers. Deeper networks enable the model to learn complex, hierarchical features in multiple stages.
- **Width:** The number of neurons in each layer. Wider layers provide the computational "space" needed to process high-dimensional information.

The central challenge in network design is managing **Model Capacity**: building a network that is complex enough to learn the patterns in the data (low bias) but simple enough not to memorize noise (low variance).

### 1.1 Theoretical Foundation: The Universal Approximator

Our exploration is framed by the **Universal Approximation Theorem** (Cybenko, 1989), which theoretically states that a single-layer network, given sufficient width, can approximate any continuous function. However, practical results (demonstrated below) reveal that adding **Depth** is often a more parameter-efficient and stable approach for handling the vast feature spaces found in real-world data.

## 2. Experimental Setup and Methodology


### 2.1 The Dataset

I used a synthesized, multi-feature classification dataset (`sklearn.datasets.make_classification`) to ensure our experiments are fully reproducible and the complexity of the problem is controlled.

**Methodological Note:** Feature scaling using `StandardScaler` was applied to ensure all features had a uniform magnitude. This is crucial for MLPs, as it prevents features with large values from dominating the gradient descent process.

## 2.2 Creative Tool: Interactive Capacity Demonstration

To provide an immediate, hands-on understanding of architectural impact, I implemented an interactive tool using `ipywidgets`. This allows the user to dynamically adjust the Depth and Width of a live network and observe the resulting changes in test accuracy and training time.



**Key Observation:** The interactive tool visually confirms that increasing capacity quickly leads to diminishing returns in accuracy, alongside a rapid increase in computational expense.


### 3. Exhaustive Analysis: The Grid Search

To systematically find the optimal configuration, I conducted an exhaustive **Grid Search** across a defined space: 1 to 5 layers (Depth) and 8 to 256 neurons (Width).

I utilized **parallel processing** (`joblib.Parallel`) to efficiently train all 25 network configurations simultaneously, demonstrating a professional approach to handling computationally intensive ML experiments.

#### 3.1 The Performance Landscape

The results of the grid search are visualized on a 3D surface plot, illustrating the total performance landscape as a function of both parameters.




##### Interpretation:

- The plot identifies a "**Goldilocks Zone**" of optimal generalization where test accuracy peaks.
- This zone is concentrated in the **moderate capacity range** (typically 2-4 layers and 32-128 neurons), reinforcing that over-capacity is unnecessary and inefficient.

#### 3.2 Analysis of Training Dynamics and Bias-Variance


I also analyzed how capacity impacts the difference between training accuracy and test accuracy, which reveals the **Bias–Variance Trade-off**.



This visualization shows that maximum capacity models achieve near-perfect training accuracy but suffer a significant drop in test accuracy, which is the classic signature of **High Variance (Overfitting)**.

## 4. Visual Proof: The Decision Boundary

To demonstrate *why* capacity matters, I trained three models of varying sizes on a simplified 2D dataset to visualize the **decision boundary**—the line the model uses to separate classes.



**Interpretation:** This visualization provides the most compelling evidence of the trade-off. The small network's inability to conform to the data's complexity (High Bias) contrasts sharply with the massive network's over-fitting to individual noise points (High Variance).

## 5. Conclusion: Guiding Architectural Choices

This tutorial has demonstrated that the secret to a successful MLP lies not in building the largest network possible, but in selecting an architecture that perfectly matches the complexity of the problem.

### Key Takeaways for Practice:

- **Prioritize Balance:** The best performance occurs in networks that exhibit **moderate capacity**, avoiding the extremes of high bias (underfitting) and high variance (overfitting).
- **Depth for Efficiency:** While the Universal Approximation Theorem focuses on width, our empirical analysis supports the preference for **moderate depth** (e.g., 2–4 layers) as a more parameter-efficient method for building feature hierarchies.
- **Reproducibility:** The complete, runnable code for all experiments is provided in the accompanying Jupyter Notebook, including professional-grade optimization techniques like parallel processing.

## References

1. Cybenko, G. (1989). "Approximation by superpositions of a sigmoidal function". *Mathematics of Control, Signals and Systems*, 2(4), 303-314.
2. Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.
3. Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
4. Scikit-Learn Documentation: *Neural network models (supervised)*. Available at: [https://scikit-learn.org/stable/modules/neural\\_networks\\_supervised.html](https://scikit-learn.org/stable/modules/neural_networks_supervised.html)

### Github Repository:

<https://github.com/MuhammadAakash/ML-tutorial-MLP-Depth-and-Width>